

# AIOps Concepts and Components: [Your Project/Repo Name]

## 1. Introduction: LLMOps with Prompt Flow

 alt text

- **What is LLMOps?** According to the AI Engineer's Handbook, LLMOps involves "deploying, monitoring, and maintaining LLMs in production; encompassing model versioning, deployment, monitoring, and maintenance best practices."
- **What is the objective?** This project, we aim to provides an LLMOps (AIOps for Large Language Models) template and guidance for building, experimenting, evaluating, and deploying LLM-infused applications using Prompt Flow.
- **Core Idea:** Bringing engineering rigor and automation to the lifecycle of LLM applications.

# 1.1 Presentation Agenda

- **Introduction & Overview:** Understanding LLMOps and its importance
- **Challenges & Solutions:** Problems addressed by AIOps template
- **Core Concepts & Components:** Building blocks of the AIOps framework
- **Architecture Design:** Technical implementation details
- **Implementation Workflow:** End-to-end process from development to deployment
- **Use Case Applications:** QCP and CMC implementation examples
- **Best Practices & Future Directions:** Guidelines and roadmap
- **Q&A Session:** Discussion and clarification

## 1.2 Supported Platforms & Execution

- **Supported Platforms:**

- Azure AI Studio
- Azure Machine Learning (AML)

- **Execution Flexibility:**

- Local execution for development and testing.
- Azure-based execution for scalability and production.

## 1.2 Flow Types & Orchestration

- **Supported Flow Types:**
  - Flexible Flows: Python Function-based, Python Class-based.
  - DAG Flows: YAML-based.
  - Automatic detection and execution of flow types.
- **CI/CD Orchestration:**
  - GitHub Actions
  - Azure DevOps
  - Jenkins
- **Focus:**
  - Inner-Loop: Experimentation and Evaluation.
  - Outer-Loop: Deployment and Inferencing.

## 2. Challenges in LLMOps Addressed

- **Managing Complexity:** Handling multiple LLM flows, each with unique lifecycles from experimentation to production.
- **Experimentation Rigor:** Systematically managing prompt variants, hyperparameters, and evaluating their performance.
- **Deployment Consistency:** Ensuring smooth and reliable deployments across different environments.
- **Data Management:** Bringing discipline to data preparation for training, experimentation, and evaluation (DataOps).
- **Reducing Boilerplate:** Enabling configuration-driven development to focus on core logic.

### 3. Core AIOps/LLMOps Concepts in this Repository

- **Centralized Code Hosting:** A single repository structure to manage multiple Prompt Flow use cases.
- **Independent Lifecycle Management:** Each flow (use case) has its own lifecycle from local development to production.
- **Variant and Hyperparameter Experimentation:** Robust support for defining and evaluating multiple configurations for flows.

## 3.1 More Core AIOps/LLMOps Concepts

- **A/B Deployment:** Facilitates comparing different flow versions in real-world settings.
- **Many-to-Many Dataset/Flow Relationships:** Allows using multiple datasets for each standard and evaluation flow.
- **Multiple Deployment Targets:** Configuration-driven deployment to:
  - Kubernetes (including ARC-enabled)
  - Azure Web Apps
  - Azure ML/AI Studio Managed Compute

## 3.2 Additional Core AIOps/LLMOps Concepts

- **Comprehensive Reporting:** Automated generation of CSV and HTML reports for experiment runs and evaluation metrics.
- **Configuration-Based Development:** Minimizing custom code through declarative configurations (e.g., `experiment.yaml`, `deployment_config.json`).
- **DataOps Integration:** Separating data pipelines from prompt engineering flows, managing datasets as versioned assets in Azure ML.



## 4. Key Components & Features

- **Prompt Flow:** The core engine for developing, evaluating, and deploying LLM workflows.
- **Standardized Folder Structure:**
  - `.azure-pipelines/` , `.github/` , `.jenkins/` : CI/CD pipeline definitions.
  - `configs/` : Deployment configurations ( `deployment_config.json` ).
  - `data/` : Raw data files for flows (e.g., `.jsonl` ).
  - `environment/` : Dockerfiles and environment specifications ( `env.yaml` ).

## 4.1 More Key Components & Features

- **Standardized Folder Structure (continued):**
  - `flows/` : Contains standard and evaluation prompt flows.
  - `tests/` : Unit tests for flows.
  - `data-pipelines/` (Optional): For DataOps implementation.
  - `llmops/` : Core Python modules for flow execution, evaluation, deployment.
  - `dataops/` : Core Python modules for DataOps pipelines.

## 4.2 Configuration Files

- `experiment.yaml` :
  - Central configuration file for each use case.
  - Defines flow paths, connections (e.g., to Azure OpenAI), datasets (sources, mappings), and evaluators.
  - Supports environment-specific overlays (e.g., `experiment.dev.yaml` , `experiment.pr.yaml` ).
- `config.py` (in `llmops/`):
  - Global setting ( `EXECUTION_TYPE` ) to switch between `LOCAL` and `AZURE` execution.

## 4.3 CI/CD & Secrets Management

- **CI/CD Automation:**

- **PR Validation:** Automated checks on Pull Requests (linting, unit tests, run on minimal data).
- **CI Pipelines:** Triggered on merges to main/development branches for full build, test, evaluation, and deployment.
- **Steps:** Registering data assets, running bulk experiments, executing evaluation flows, deploying endpoints, testing endpoints.

- **Secrets Management:**

- **Local:** `.env` file at the root (gitignored).
- **Cloud (GitHub):** `ENV_VARS` repository secret.
- **Cloud (Azure DevOps):** Library variable groups.
- Placeholders like `${SECRET_NAME}` used in configuration files.

## 4.4 Data Management & Example Use Cases

- **Data Management (DataOps):**
  - `dataops_config.json`: Configuration for data pipelines.
  - Scripts to process raw data and register it as Azure ML Data Assets.
  - Flows consume data from these registered assets.
- **Example Use Cases Provided:**
  - Web Classification (YAML-based)
  - Named Entity Recognition (YAML-based)
  - Math Coding (YAML-based)
  - Chat with PDF (RAG-based, YAML)
  - Code Generation (Function-based)
  - Chat Application (Class-based)

## 5. High-Level Workflow

### 1. Local Development & Experimentation:

- Define/modify flows (standard and evaluation).
- Configure `experiment.yaml` for the use case.
- Set up `.env` for local secrets.
- Run experiments and evaluations locally using provided Python scripts (`llmops.common.prompt_pipeline`, `llmops.common.prompt_eval`).

### 2. Source Control & PR Validation:

- Commit changes to a feature branch.
- Create a Pull Request.
- Automated PR pipeline runs (linting, tests, minimal flow execution).

## 5.1 CI/CD Pipeline & Post-Deployment

### 3. CI/CD Pipeline (on Merge):

- **Setup:** Authenticate to Azure, install dependencies.
- **Data Registration:** Register/update datasets in Azure ML.
- **Bulk Run (Experimentation):** Execute standard flow(s) with variants against datasets on Azure.
- **Evaluation:** Execute evaluation flow(s) using the outputs of the bulk run.
- **Reporting:** Generate and publish metrics reports.
- **(Optional) Manual Approval Gate:** Human validation of metrics.
- **Deployment:**
  - Build Docker image (if deploying to Web Apps/AKS).
  - Deploy flow to the configured target (AML Managed Endpoint, AKS, Web App).

## 6. Benefits for Your Client

- **Accelerated Development:** Faster iteration on LLM features due to streamlined processes and automation.
- **Improved Quality & Reliability:** Rigorous testing, evaluation, and consistent deployments.
- **Scalability:** Easily scale LLM applications using Azure's cloud infrastructure.
- **Cost Efficiency:** Optimized resource usage through managed compute and efficient experimentation.



## 6.1 More Benefits for Your Client

- **Enhanced Collaboration:** Standardized tools and processes for data scientists, ML engineers, and DevOps.
- **Reproducibility:** Versioned code, data, and configurations ensure experiments and deployments are reproducible.
- **Flexibility:** Supports various flow types and deployment targets to fit diverse needs.

## 7. Q&A

## 8. AIOps Design Architecture

- **Architecture Overview:** A comprehensive LLMOps architecture supporting both experimentation and deployment of large language model applications.
- **Development Environment:**
  - Local execution capabilities with Python scripts
  - VS Code integration through extensions
  - Azure AI Studio-compatible workflows
- **Flow Types & Implementation:**
  - YAML-based flows: Traditional DAG-based workflows (e.g., web\_classification)
  - Function-based flows: Python functions with prompt flow capabilities
  - Class-based flows: Python classes with more complex state management

## 8.1 AIOps Design - Key Components

- **Experimentation & Evaluation (Inner Loop):**
  - Variant testing with multiple prompt configurations
  - Comprehensive metrics collection and reporting
  - Multi-dataset support for robust evaluation
  - A/B testing capabilities
- **Deployment Targets (Outer Loop):**
  - Azure ML Compute with managed endpoints
  - Kubernetes deployments (including AKS)
  - Azure Web Apps using Docker containers
  - Support for A/B deployment strategies
- **DataOps Integration:**

## 8.2 AIOps Design - Architecture Benefits

- **Complete Lifecycle Management:** End-to-end coverage from experimentation to production
- **Flexibility in Development:** Multiple flow types to suit various use cases
- **Robust Evaluation:** Comprehensive evaluation with multiple metrics and datasets
- **Deployment Options:** Multiple Azure-based deployment targets to meet different requirements:
  - Scalability with Azure Kubernetes Service
  - Managed services with Azure ML endpoints
  - Cost-effective options with Azure Web Apps
- **CI/CD Integration:** Automated workflows ensuring quality and accelerating delivery

## 9. Use Case Application: QCP and CMC

- **Quality Control Platform (QCP):**

- Streamlines generation of quality control reports by integrating disparate data sources
- Uses LLMs to analyze data and generate reports based on predefined rules
- Currently in development stage with Azure DevOps deployment

- **CMC (Content Management and Compliance):**

- RAG-based solution for finding relevant historical documents and generating summaries
- Helps users find similar questions/answers from past documents during submission processes
- Uses Azure AI Search and OpenAI models from the AI marketplace

## 9.1 QCP and CMC Requirements

- **QCP Requirements:**

- Data processing pipeline for bacteria samples analysis
- Manual data ingestion with temporary processing (no permanent storage)
- High UI availability with low latency for real-time analysis
- Content moderation services for model safety
- Evaluation using golden datasets and user feedback

- **CMC Requirements:**

- Timer-triggered data ingestion pipeline (daily/weekly)
- Vector search capabilities for document retrieval
- Confidential data handling within Novo network
- Document summarization capabilities

## 9.2 How AIOps Addresses These Requirements

- **For QCP:**

- Standardized prompt flow development for data analysis logic
- Evaluation flows for measuring accuracy with golden datasets
- Automated deployment pipelines through Azure DevOps integration
- Monitoring capabilities for performance tracking
- Content filtering and prompt shielding capabilities

- **For CMC:**

- RAG-based flow templates (e.g., Chat with PDF example)
- Azure AI integration for vector search functionality
- Configurable data pipeline structure through DataOps modules
- Secure deployment within Azure environments



## 9.3 AIOps Limitations and Future Enhancements

- **Current Limitations:**

- Limited support for Databricks integration in data pipelines
- Need for custom connectors to Azure AI Search beyond standard templates
- Lack of specialized metrics for RAG evaluation in the CMC use case
- Limited fine-tuning options for both applications' specific domains
- No built-in content moderation services (relies on Azure's services)

- **Recommended Enhancements:**

- Develop Databricks-specific DataOps connectors for CMC
- Implement specialized RAG evaluation metrics for search quality
- Add domain-specific prompt templates for QCP's bacteria analysis
- Enhance monitoring dashboards for usability metrics

## 10. LLMOps Components in Detail

- **Definition:** According to the AI Engineer's Handbook, "LLMOps is a set of practices and tools for deploying, monitoring, and maintaining Large Language Models in production. It extends MLOps principles specifically for LLM applications."

 Development to Production Workflow for LLMs

# 10.1 LLMOps Key Components - Deployment

- **Model Versioning and Deployment:**
  - Managing different model versions
  - Enabling seamless rollbacks and updates
  - Supporting A/B testing for model variants
- **Infrastructure Management:**
  - Setting up necessary hardware and software environments
  - Ensuring efficient model execution
  - Configuring deployment targets (AKS, Azure ML, Web Apps)
- **Scaling and Performance Optimization:**
  - Adjusting resources based on demand
  - Implementing horizontal and vertical scaling strategies

## 10.2 LLMOps Key Components - Monitoring

- **Response Quality Tracking:**
  - Assessing generated response quality
  - Validating against user expectations
- **Performance Metrics:**
  - Measuring latency and throughput
  - Evaluating model efficiency in real-time
- **Usage Analytics:**
  - Analyzing user interaction patterns
  - Identifying improvement opportunities
- **Error Monitoring:**

## 10.3 LLMOps Key Components - Maintenance

- **Model Updates and Versioning:**
  - Incorporating new data and improvements
  - Maintaining model relevancy and effectiveness
- **Data Pipeline Management:**
  - Overseeing data flow into and out of models
  - Ensuring clean, relevant, and timely data
- **Fine-tuning Workflows:**
  - Adjusting model parameters based on feedback
  - Retraining with new datasets
- **Security Patches:**

## 10.4 LLMOps Implementation Steps

- **Step 1: Select a Foundation Model**
  - Choose between proprietary models (OpenAI's GPT) or open-source options
  - Consider performance requirements vs. cost constraints
- **Step 2: Adapt to Downstream Tasks**
  - Implement prompt engineering techniques
  - Apply fine-tuning for specific applications
  - Incorporate external data through RAG
  - Utilize embeddings for search and recommendations
- **Step 3: Deploy and Maintain**
  - Establish version control and governance
  - Implement monitoring and feedback loops

# 11. LLMOps Security Best Practices

- **Access Control Implementation:**

- Role-based access control for model APIs
- Multi-factor authentication for sensitive operations
- Strict permission boundaries between environments

- **Data Privacy Protections:**

- Data encryption at rest and in transit
- Personally Identifiable Information (PII) detection and redaction
- Compliance with regulatory frameworks (GDPR, HIPAA, etc.)

- **Prompt Injection Defenses:**

- Input validation and sanitization

◦ Context boundaries enforcement

## 12. LLMOps Development Best Practices

- **Version Control for Prompts:**

- Maintain all prompts in source control
- Document prompt changes with detailed commit messages
- Track prompt performance across versions

- **Testing Frameworks:**

- Implement automated testing for prompts and flows
- Create comprehensive test suites for different scenarios
- Test for edge cases and potential vulnerabilities

- **CI/CD Implementation:**

- Automate flow deployment with quality gates

Implement progressive deployment strategies



# 13. LLMOps Production Best Practices

- **Load Balancing Strategies:**
  - Distribute traffic across multiple model instances
  - Implement auto-scaling based on demand patterns
  - Optimize resource allocation for cost efficiency
- **Failover Mechanisms:**
  - Design redundant systems for high availability
  - Implement graceful degradation patterns
  - Create fallback responses for service interruptions
- **Caching Implementation:**
  - Cache common responses to reduce latency
  - Implement tiered caching strategies

## 14. Azure-Specific Implementation Guidance

- **Azure OpenAI Integration:**
  - Leverage Azure OpenAI Service for compliance and security
  - Implement Managed Identity for secure authentication
  - Use Private Endpoints for network isolation
- **Deployment Patterns:**
  - Blue-Green deployments for zero-downtime updates
  - Canary releases for controlled feature rollout
  - Shadow deployments for performance testing
- **Monitoring Setup:**
  - Azure Monitor and Application Insights integration
  - Custom dashboards for LLM-specific metrics

# 15. Summary and Next Steps

- **Key Takeaways:**

- LLMOps brings engineering rigor to AI development
- Our AIOps template provides complete lifecycle management
- Standardized approach improves quality and accelerates delivery

- **Getting Started:**

- Clone the template repository
- Follow the documentation for your specific use case
- Start with the simplest flow type for your needs

- **Future Roadmap:**

- Enhanced RAG metrics and evaluation capabilities

Additional connectors for data sources

## 16. Q&A

- Questions?
- **Demo Available Upon Request**
- **Contact:** [Your Contact Information]

